

29702361—SQL Programming & Creative Writing

Task 1

- Considering the attached `Student` and `Student_Courses` relations, discuss when Union, Inner Join, Left Outer Join, Right Outer Join, and Full Outer Join should be used and provide examples using the attached relations.

Cite any source in APA format.

UNION

The `UNION` operations enables queries to combine their outputs if the provided attributes contain similar data types and attribute names (Elmasri & Navathe, 2016). A case in point is where you want to list the `Student_ID` attributes from both `Student` and `Student_Courses`. Using MySQL, for instance, that could be achieved by executing the query:

```
1 | SELECT Student_ID FROM Student
2 | UNION
3 | SELECT Student_ID FROM Student_Courses;
```

This would produce the output:

```
1 | +-----+
2 | | Student_ID |
3 | +-----+
4 | | 0000      |
5 | | 1111      |
6 | | 2222      |
7 | | 3333      |
8 | | 4444      |
9 | | 5555      |
10 | +-----+
11 | 6 rows in set (0.001 sec)
```

This lists the distinct rows containing the `Student_ID` values. On the other hand, by tweaking the earlier query to use the `UNION ALL` operation we get the output:

```
1 | +-----+
2 | | Student_ID |
3 | +-----+
4 | | 0000      |
5 | | 1111      |
6 | | 2222      |
7 | | 3333      |
8 | | 4444      |
9 | | 5555      |
10 | | 2222      |
11 | | 2222      |
12 | | 2222      |
13 | | 0000      |
14 | | 1111      |
15 | | 1111      |
16 | | 0000      |
17 | | 1111      |
18 | | 1111      |
19 | | 3333      |
20 | | 3333      |
21 | | 4444      |
22 | | 4444      |
23 | | 4444      |
24 | +-----+
25 | 20 rows in set (0.001 sec)
```

INNER JOIN

The `INNER JOIN` operation lists all the rows from the specified tables that meet the a given criteria (Elmasri & Navathe, 2016). As an example, if you want to select all the rows from with `Student_Name`, `Dept_Name`, and `Course_ID` attributes from the `Student` and `Student_Courses` relations; where, the `Student_ID` values match in both relations and the `Grade` attribute in `Student_Courses` has an `A` value; you would write such a query as:

```
1 | SELECT Student_Name, Dept_Name, Course_ID
2 | FROM Student
3 | INNER JOIN Student_Courses
4 | ON Student.Student_ID = Student_Courses.Student_ID
5 | AND Student_Courses.Grade = 'A';
```

Which would then produce the output of:

1	+-----+-----+-----+		
2	Student_Name Dept_Name Course_ID		
3	+-----+-----+-----+		
4	AAAAA Computer Science CS 1101		
5	AAAAA Computer Science CS 2203		
6	BBBBB Computer Science CS 2204		
7	BBBBB Computer Science CS 2401		
8	DDDDD Education EDUC 5010		
9	EEEEE Physics PHY 1101		
10	EEEEE Physics PHY 3304		
11	+-----+-----+-----+		
12	7 rows in set (0.001 sec)		

LEFT OUTER JOIN

The `LEFT OUTER JOIN` operation "returns not only the rows matching the join condition (that is, rows with matching values in the common column), but also the rows in the left side table with unmatched values in the right side table" (Rob & Coronel, 2009 , pg. 310). As an example, say you want to list all the available `Student` relations using their `Student_ID` and `Student_Name` attributes together with the courses that they do, you would write a query such as:

```
1 | SELECT Student.Student_ID, Student_Name, Course_ID
2 | FROM Student
3 | LEFT OUTER JOIN Student_Courses
4 | ON Student.Student_ID = Student_Courses.Student_ID;
```

On execution, you will notice that the `Student` named `FFFFF` has not enrolled in any course:

1	+-----+-----+-----+		
2	Student_ID Student_Name Course_ID		
3	+-----+-----+-----+		
4	0000 AAAAA CS 1101		
5	0000 AAAAA CS 2203		
6	1111 BBBBB CS 1101		
7	1111 BBBBB CS 1102		
8	1111 BBBBB CS 2204		
9	1111 BBBBB CS 2401		
10	2222 CCCCC BUS 1101		
11	2222 CCCCC BUS 2201		
12	2222 CCCCC BUS 3302		
13	3333 DDDDD EDUC 5010		
14	3333 DDDDD EDUC 5210		
15	4444 EEEEE PHY 1101		
16	4444 EEEEE PHY 2202		
17	4444 EEEEE PHY 3304		
18	5555 FFFFF NULL		
19	+-----+-----+-----+		
20	15 rows in set (0.001 sec)		

Additionally, because the left table in this case is the `Student` relation, the ordering of the `Student_ID` attributes in the `Student` relation define how the rest of the columns are sorted.

RIGHT OUTER JOIN

The `RIGHT OUTER JOIN` operation "returns not only the rows matching the join condition (that is, rows with matching values in the common column), but also the rows in the right side table with unmatched values in the left side table" (Rob & Coronel, 2009 , pg. 310). As an example, say you want to list all only those courses that students have took in a table containing the `Student_ID`, `Student_Name`, and `Course_ID` relations, you would write a query such as:

```
1 SELECT Student.Student_ID, Student_Name, Course_ID
2 FROM Student
3 RIGHT OUTER JOIN Student_Courses
4 ON Student.Student_ID = Student_Courses.Student_ID;
```

Also, this output shows how one can eliminate the `NULL` values from the previous section:

1	+-----+-----+-----+
2	Student_ID Student_Name Course_ID
3	+-----+-----+-----+
4	2222 CCCCC BUS 1101
5	2222 CCCCC BUS 2201
6	2222 CCCCC BUS 3302
7	0000 AAAAA CS 1101
8	1111 BBBBB CS 1101
9	1111 BBBBB CS 1102
10	0000 AAAAA CS 2203
11	1111 BBBBB CS 2204
12	1111 BBBBB CS 2401
13	3333 DDDDD EDUC 5010
14	3333 DDDDD EDUC 5210
15	4444 EEEEE PHY 1101
16	4444 EEEEE PHY 2202
17	4444 EEEEE PHY 3304
18	+-----+-----+-----+
19	14 rows in set (0.001 sec)

FULL OUTER JOIN

The `FULL OUTER JOIN` operation "returns not only the rows matching the join condition (that is, rows with matching values in the common column), but also all of the rows with unmatched values in either side table" (Rob & Coronel, 2009 , pg. 311). In MySQL, a `FULL OUTER JOIN` operation can be realized by combining the outputs of the `LEFT OUTER JOIN` and `RIGHT OUTER JOIN` operations using a `UNION` operator—as this query shows:

```
1 SELECT Student.Student_ID, Student_Name, Course_ID
2 FROM Student
3 LEFT OUTER JOIN Student_Courses
4 ON Student.Student_ID = Student_Courses.Student_ID
5 UNION
6 SELECT Student.Student_ID, Student_Name, Course_ID
7 FROM Student
8 RIGHT OUTER JOIN Student_Courses
9 ON Student.Student_ID = Student_Courses.Student_ID;
```

The resultant output:

1	+-----+-----+-----+
2	Student_ID Student_Name Course_ID
3	+-----+-----+-----+
4	0000 AAAAA CS 1101
5	0000 AAAAA CS 2203
6	1111 BBBBB CS 1101
7	1111 BBBBB CS 1102
8	1111 BBBBB CS 2204
9	1111 BBBBB CS 2401
10	2222 CCCCC BUS 1101
11	2222 CCCCC BUS 2201
12	2222 CCCCC BUS 3302
13	3333 DDDDD EDUC 5010
14	3333 DDDDD EDUC 5210

15		4444		EEEE		PHY 1101	
16		4444		EEEE		PHY 2202	
17		4444		EEEE		PHY 3304	
18		5555		FFFF		NULL	
19	+-----+						
20	15 rows in set (0.001 sec)						

References

Elmasri, R., & Navathe, S. (2016). *Fundamentals of database systems*. Pearson.

Rob, P., & Coronel, C. (2009). *Database systems: Design, implementation, and management*. Thomson/Course Technology.

Task 2

- Create two different “Views” using the attached relations, then Alter one of the Views and Drop the other one. After that, discuss why Views are used by giving examples using the attached relations.
- Describe what you did (This does not mean that you copy and paste from what you have posted or the assignments you have prepared. You need to describe what you did and how you did it), what you learned, your weekly activities, in what ways are you able to apply the ideas and concepts gained, and finally, describe one important thing that you are thinking about in relation to the activity.

Cite any source in APA format.

In a case where one wants to be able to list the details of a `Student` relation using only the `Student_ID` and `Student_Name` attributes, it would be convenient to create a View that executes a given `SELECT` on the underlying table instead of having to run that query repeatedly by oneself. Thus to create such a View (named `student_details`), one would create it using the DDL command:

```
1 CREATE VIEW student_details AS
2     SELECT Student_ID, Student_Name
3     FROM Student;
```

Accordingly, to use the View, it is as simple as running the query:

```
1 SELECT * FROM student_details;
```

Which produces the output:

1		+-----+					
2		Student_ID		Student_Name			
3		+-----+					
4		0000		AAAAA			
5		1111		BBBBB			
6		2222		CCCCC			
7		3333		DDDDD			
8		4444		EEEE			
9		5555		FFFFF			
10		+-----+					
11		6 rows in set (0.001 sec)					

Yet, this View does not offer enough detail when one wants to know what marks a student scored in a course, example. Thus, it is necessary to create another View named `student_marks`, which contains both the identifying details of the students and their course grades:

```
1 CREATE VIEW student_marks AS
2     SELECT Student_Name, Dept_Name, Course_ID, Year, Grade FROM Student
3     LEFT OUTER JOIN Student_Courses ON
4     Student.Student_ID = Student_Courses.Student_ID;
```

On querying the `student_marks` View:

```
1 SELECT * FROM student_marks;
```

The resulting output is:

--	--	--	--	--	--	--	--

1	+-----+-----+-----+-----+-----+										
2		Student_Name		Dept_Name		Course_ID		Year		Grade	
3	+-----+-----+-----+-----+-----+										
4		AAAAA		Computer Science		CS 1101		2019		A	
5		AAAAA		Computer Science		CS 2203		2020		A	
6		BBBBB		Computer Science		CS 1101		2019		B	
7		BBBBB		Computer Science		CS 1102		2019		B	
8		BBBBB		Computer Science		CS 2204		2020		A	
9		BBBBB		Computer Science		CS 2401		2020		A	
10		CCCCC		Business Administration		BUS 1101		2019		C	
11		CCCCC		Business Administration		BUS 2201		2020		B	
12		CCCCC		Business Administration		BUS 3302		2020		B	
13		DDDDD		Education		EDUC 5010		2019		A	
14		DDDDD		Education		EDUC 5210		2020		B	
15		EEEEEE		Physics		PHY 1101		2019		A	
16		EEEEEE		Physics		PHY 2202		2020		C	
17		EEEEEE		Physics		PHY 3304		2020		A	
18		FFFFF		History		NULL		NULL		NULL	
19	+-----+-----+-----+-----+-----+										
20	15 rows in set (0.001 sec)										

Still, the `student_marks` View is deficient in one aspect. It does not list the `Student_ID` attributes. Hence, it requires some modification, which can be achieved by altering it using the following DDL command:

```
1 ALTER VIEW student_marks AS
2 SELECT Student.Student_ID, Student_Name, Dept_Name, Course_ID, Year, Grade FROM Student
3 LEFT OUTER JOIN Student_Courses ON
4 Student.Student_ID = Student_Courses.Student_ID;
```

On querying the altered `student_marks` View:

```
1 SELECT * FROM student_marks;
```

The resulting output shows that a new column named `Student_ID` has been created:

1	+-----+-----+-----+-----+-----+												
2		Student_ID		Student_Name		Dept_Name		Course_ID		Year		Grade	
3	+-----+-----+-----+-----+-----+												
4		0000		AAAAA		Computer Science		CS 1101		2019		A	
5		0000		AAAAA		Computer Science		CS 2203		2020		A	
6		1111		BBBBB		Computer Science		CS 1101		2019		B	
7		1111		BBBBB		Computer Science		CS 1102		2019		B	
8		1111		BBBBB		Computer Science		CS 2204		2020		A	
9		1111		BBBBB		Computer Science		CS 2401		2020		A	
10		2222		CCCCC		Business Administration		BUS 1101		2019		C	
11		2222		CCCCC		Business Administration		BUS 2201		2020		B	
12		2222		CCCCC		Business Administration		BUS 3302		2020		B	
13		3333		DDDDD		Education		EDUC 5010		2019		A	
14		3333		DDDDD		Education		EDUC 5210		2020		B	
15		4444		EEEEEE		Physics		PHY 1101		2019		A	
16		4444		EEEEEE		Physics		PHY 2202		2020		C	
17		4444		EEEEEE		Physics		PHY 3304		2020		A	
18		5555		FFFFF		History		NULL		NULL		NULL	
19	+-----+-----+-----+-----+-----+												
20	15 rows in set (0.001 sec)												

Finally, because this `student_marks` View contains all the info we need on a student's performance—including the identifying details, the `student_details` View is rendered redundant and removing it would be welcome. This is attainable using the `DROP VIEW` operation:

```
1 DROP VIEW student_details;
```

As a result, any further use of that View identifier is bound to throw an error.

The MySQL installation contains a command line tool , which I used extensively when creating, altering, and dropping the views for this task. Also, I used it for executing `SELECT` statements that allowed me to inspect the contents of the views and their underlying tables. Thus when creating the `student_details` view, it was as straightforward as running its DDL command on the terminal. Yet, this was only the execution part of the task. First, I had to assess the

provided relations to determine which attributes would be most helpful for a user who possesses the data of students and their course grades. Second, to complete the task, I had to experiment with various types of views to come up with those that combined the most insightful data from both the provided relations. All in all, the testing would not have been complete if I had not created actual database tables containing the attributes and values displayed in the provided relation screen-shots.

All these activities reminded me how important it was to know how to generate DDL statements that could generate tables and data matching given use cases. Thus, I was pleased by how adept I have become at handling SQL tasks because I could borrow easily from skills learned in earlier SQL assignments. On the other hand, I learned that views only represent a virtual form of the tables that a database system contains. It was actually a relieving surprise to find out that I could create, alter, and even drop views without affecting the data that was held by the tables in the database. This was especially crucial because at some point I had hesitated to implement the `DROP VIEW` routines out of fear that I would delete all the data that I had entered into the tables manually.